

Problem Set 4: Neural Networks

This assignment requires a working IPython Notebook installation, which you should already have. If not, please refer to the instructions in the previous problem sets.

In part 2 (programming) of this assignment, you DO NOT need to make any modification code in this IPython Notebook. Instead you will implement your own simple neural network in the `mlp.py` file. To submit your answers insert `mlp.py` into this IPython Notebook along with the written answers and submit everything as a single PDF file.

Total: 100 points.

[30pts] Problem 1: Backprop in a simple MLP

This problem asks you to derive all the steps of the backpropagation algorithm for a simple classification network. Consider a fully-connected neural network, also known as a multi-layer perceptron (MLP), with a single hidden layer and a one-node output layer. The hidden and output nodes use an elementwise sigmoid activation function and the loss layer uses cross-entropy loss:

$$f(z) = \frac{1}{1 + \exp(-z)}$$

$$L(\hat{y}, y) = -y \ln(\hat{y}) - (1 - y) \ln(1 - \hat{y})$$

The computation graph for an example network is shown below. Note that it has an equal number of nodes in the input and hidden layer (3 each), but, in general, they need not be equal. Also, to make the application of

backprop easier, we show the *computation graph* which shows the dot product and activation functions as their own nodes, rather than the usual graph showing a single node for both.



The forward and backward computation for a single example are given below. Note that \odot stands for element-wise multiplication.

The forward step is:



and the backward step is:



Write down each step of the backward pass explicitly for all layers, i.e. compute all intermediate gradients of J above, expressing them as a function of variables x (a vector of inputs), y (a scalar label), $h^{(k)}$ (layer k outputs), $W^{(k)}$ (parameters in layer k), and $b^{(k)}$ (bias parameters of layer k). Notice that $h^{(k)}$ and $b^{(k)}$ can be scalars or vectors depending on layer k, and $W^{(k)}$ can be a matrix or a vector. NOTE: We will assume no regularization, so you can omit the terms involving Ω .

As an example, we will derive the first gradient for you, i.e. the gradient of the cross entropy loss w.r.t. its (scalar) input $\nabla_{\hat{y}} J = \nabla_{\hat{y}} L(\hat{y}, y)$.

$$\nabla_{\hat{y}} L(\hat{y}, y) = \nabla_{\hat{y}} [-y \ln(\hat{y}) - (1 - y) \ln(1 - \hat{y})] = \frac{\hat{y} - y}{(1 - \hat{y})\hat{y}} = \frac{h^{(2)} - y}{(1 - h^{(2)})h^{(2)}}$$

Please derive the remaining gradients, listed below.

Hint: you should substitute the updated values for the gradient g in each step and simplify as much as possible.

Hint: Useful information about vectorized chain rule and backpropagation:

If you are struggling with computing the vectorized version of chain rule you may find this example helpful:

<https://web.stanford.edu/class/cs224n/readings/gradient-notes.pdf>

It also contains some helpful shortcuts for computing gradients.

[5pts] Q1.1: Derive $\nabla_{a^{(2)}} J$, where $a^{(2)}$ the (scalar) pre-nonlinearity activation of layer 2.

Hint: to get this scalar value, multiply the partial value g , which should be equal to the $\nabla_{\hat{y}} L(\hat{y}, y)$ that we computed above, with the sigmoid derivative $f'(a^{(2)})$, and note that $f(a^{(2)}) = h^{(2)}$.

$$f(a^{(2)}) = h^{(2)} = \frac{1}{1 + e^{-a^{(2)}}}$$

Thus,

$$\frac{d}{da^{(2)}} \frac{1}{1 + e^{-a^{(2)}}} = \frac{-e^{-a^{(2)}}}{(1 + e^{-a^{(2)}})^2} = \frac{1}{1 + e^{-a^{(2)}}} \left(1 - \frac{1}{1 + e^{-a^{(2)}}}\right) = h^{(2)}(1 - h^{(2)})$$

$$\nabla_{a^{(2)}} J = \nabla_{\hat{y}} L(\hat{y}, y) \odot h^{(2)}(1 - h^{(2)}) = h^{(2)} - y$$

[5pts] Q1.2: $\nabla_{b^{(2)}} J$

$\nabla_{b^{(2)}} J = g$ (ignore regularization)

$$\nabla_{b^{(2)}} J = \nabla_{b^{(2)}} (b^{(2)} + W^{(2)} h^{(k-1)}) \nabla_{a^{(2)}} J = 1(h^{(2)} - y) = h^{(2)} - y$$

[5pts] Q1.3: $\nabla_{W^{(2)}} J$

Hint: this should be a vector, since $W^{(2)}$ is a vector.

$$\nabla_{W^{(2)}} J = \nabla_{W^{(2)}} (b^{(2)} + W^{(2)} h^{(1)}) \nabla_{a^{(2)}} J = (h^{(2)} - y) h^{(1)T}$$

[5pts] Q1.4: $\nabla_{h^{(1)}} J$

$$\nabla_{h^{(1)}} = \nabla_{h^{(1)}}(b^{(2)} + W^{(2)}h^{(1)})\nabla_{a^{(2)}} = \nabla_{a^{(2)}}JW^{(2)T} = (h^{(2)} - y)W^{(2)T}$$

[5pts] Q1.5: $\nabla_{b^{(1)}}J, \nabla_{W^{(1)}}J$

We have $f(a^{(1)})' = h^{(1)}(1 - h^{(1)})$

$$\nabla_{a^{(1)}}J = \nabla_{h^{(1)}}J \odot \nabla_{a^{(1)}}f(a^{(1)}) = (h^{(2)} - y)W^{(2)T} \odot h^{(1)}(1 - h^{(1)})$$

Thus we have:

$$\nabla_{b^{(1)}}J = \nabla_{b^{(1)}}(b^{(1)} + W^{(1)}h^{(k-1)})\nabla_{a^{(1)}}J = \nabla_{a^{(1)}}J = (h^{(2)} - y)W^{(2)T} \odot h^{(1)}(1 - h^{(1)})$$

$$\nabla_{W^{(1)}}J = \nabla_{W^{(1)}}(b^{(1)} + W^{(1)}h^{(0)})\nabla_{a^{(1)}}J = h^{(0)}\nabla_{a^{(1)}}J = h^{(0)}[(h^{(2)} - y)W^{(2)T} \odot h^{(1)}(1 - h^{(1)})]$$

[5pts] Q1.6 Briefly, explain how the computational speed of backpropagation would be affected if it did not include a forward pass

Without forwardpass, no activation functions h are stored, speed increase significantly since we have to compute them repeatedly from the start every time we get to the node.

[50pts] Problem 2 (Programming): Implementing a simple MLP

In this problem we will develop a neural network with fully-connected layers, or Multi-Layer Perceptron (MLP). We will use it in classification tasks. This question is adapted from [Stanford CS231n](#).

In the current directory, you can find a file `m1p.py`, which contains the definition for class `TwoLayerMLP`. As the name suggests, it implements a 2-layer MLP, or MLP with 1 *hidden* layer. You will implement your code in the same file, and call the member functions in this notebook. Below is some initialization. The `autoreload` command makes sure that `m1p.py` is periodically reloaded.

In [1]:

```
# setup
import numpy as np
```

```
import matplotlib.pyplot as plt
from mlp import TwoLayerMLP

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

Next we initialize a toy model and some toy data, the task is to classify five 4-d vectors.

In [2]:

```
# Create a small net and some toy data to check your implementations.
# Note that we set the random seed for repeatable experiments.
input_size = 4
hidden_size = 10
num_classes = 3
num_inputs = 5

def init_toy_model(activ, std=1e-1):
    np.random.seed(0)
    return TwoLayerMLP(input_size, hidden_size, num_classes, std=std, activation=activ)

def init_toy_data():
    np.random.seed(1)
    X = 10 * np.random.randn(num_inputs, input_size)
    y = np.array([0, 1, 2, 2, 1])
    return X, y

X, y = init_toy_data()
```

```
print('X = ', X)
print()
print('y = ', y)
```

```
X = [[ 16.24345364 -6.11756414 -5.28171752 -10.72968622]
      [  8.65407629 -23.01538697  17.44811764 -7.61206901]
      [  3.19039096 -2.49370375  14.62107937 -20.60140709]
      [ -3.22417204 -3.84054355  11.33769442 -10.99891267]
      [ -1.72428208 -8.77858418  0.42213747  5.82815214]]
```

```
y = [0 1 2 2 1]
```

[5pts] Q2.1 Forward pass: Sigmoid

Our 2-layer MLP uses a softmax output layer (**note**: this means that you don't need to apply a sigmoid on the output) and the multiclass cross-entropy loss to perform classification.

Softmax function:

For class j :

$$P(y_{(j)}|x) = \frac{\exp(z_j)}{\sum_{c=1}^C \exp(z_c)}$$

Where C is the number of classes and z is class-wise output of the network.

Multiclass cross-entropy loss function:

$$J = \frac{1}{m} \sum_{i=1}^m \sum_{c=1}^C [-y_{(c)} \log(P(y_{(c)}|x^{(i)}))]$$

$y_{(c)} = 1$ for the ground truth class and 0 otherwise.

m is the number of inputs in a batch and C is the number of classes.

Please take a look at method `TwoLayerMLP.loss` in the file `m1p.py`. This function takes in the data and weight parameters, and computes the class scores (aka logits), the loss L , and the gradients on the parameters.

- Complete the implementation of forward pass (up to the computation of `scores`) for the sigmoid activation: $\sigma(x) = \frac{1}{1+\exp(-x)}$.

Note 1: Softmax cross entropy loss involves the [log-sum-exp operation](#). This can result in numerical underflow/overflow. Read about the solution in the link, and try to understand the calculation of `loss` in the code.

Note 2: You're strongly encouraged to implement in a vectorized way and avoid using slower `for` loops. Note that most numpy functions support vector inputs.

Check the correctness of your forward pass below. The difference should be very small ($< 1e-6$).

In [3]:

```
net = init_toy_model('sigmoid')
loss, _ = net.loss(X, y, reg=0.1)
correct_loss = 1.182248
print(loss)
print('Difference between your loss and correct loss:')
print(np.sum(np.abs(loss - correct_loss)))
```

```
1.1822479803941373
```

```
Difference between your loss and correct loss:
```

```
1.9605862711102873e-08
```

[10pts] Q2.2 Backward pass: Sigmoid

- For sigmoid activation, complete the computation of `grads`, which stores the gradient of the loss with respect to the variables `w1`, `b1`, `w2`, and `b2`.

Now debug your backward pass using a numeric gradient check. Again, the differences should be very small.

```
In [4]: # Use numeric gradient checking to check your implementation of the backward pass.
# If your implementation is correct, the difference between the numeric and
# analytic gradients should be less than 1e-8 for each of W1, W2, b1, and b2.
from utils import eval_numerical_gradient

loss, grads = net.loss(X, y, reg=0.1)

# these should all be very small
for param_name in grads:
    f = lambda W: net.loss(X, y, reg=0.1)[0]
    param_grad_num = eval_numerical_gradient(f, net.params[param_name], verbose=False)
    print('%s max relative error: %e'%(param_name, rel_error(param_grad_num, grads[param_name])))
```

```
W2 max relative error: 8.048892e-10
b2 max relative error: 5.553999e-11
W1 max relative error: 1.126755e-08
b1 max relative error: 2.035406e-06
```

[5pts] Q2.3 Train the Sigmoid network

To train the network we will use stochastic gradient descent (SGD), implemented in `TwoLayerNet.train`. Then we train a two-layer network on toy data.

- Implement the prediction function `TwoLayerNet.predict`, which is called during training to keep track of training and validation accuracy.

You should get the final training loss around 0.1, which is good, but not too great for such a toy problem. One problem is that the gradient magnitude for `W1` (the first layer weights) stays small all the time, and the neural net doesn't get much "learning signals". This has to do with the saturation problem of the sigmoid activation function.

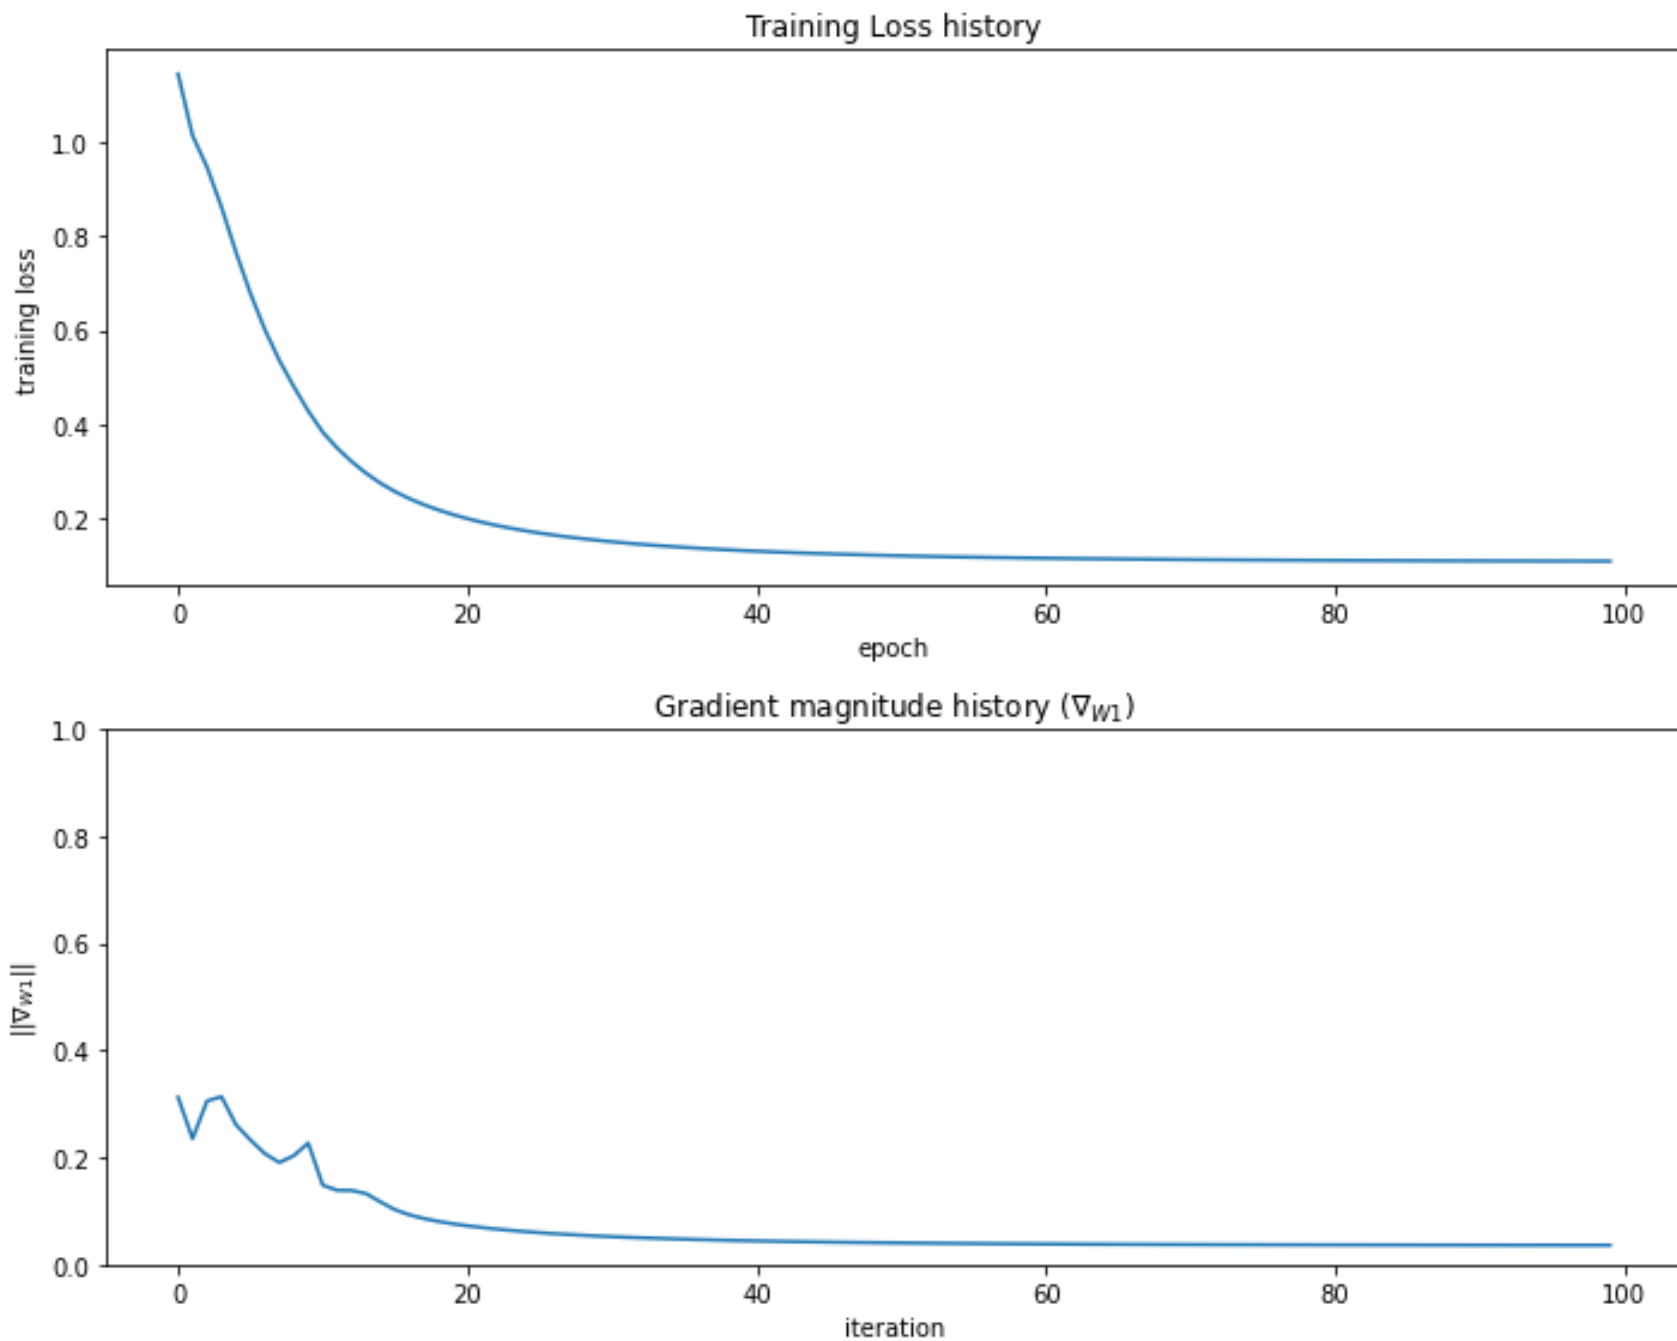
```
In [5]: net = init_toy_model('sigmoid', std=1e-1)
stats = net.train(X, y, X, y,
                  learning_rate=0.5, reg=1e-5,
```



```
num_epochs=100, verbose=False)
print('Final training loss: ', stats['loss_history'][-1])

# plot the loss history and gradient magnitudes
fig, (ax1, ax2) = plt.subplots(2, 1)
ax1.plot(stats['loss_history'])
ax1.set_xlabel('epoch')
ax1.set_ylabel('training loss')
ax1.set_title('Training Loss history')
ax2.plot(stats['grad_magnitude_history'])
ax2.set_xlabel('iteration')
ax2.set_ylabel(r'$||\nabla_{W1}||$')
ax2.set_title('Gradient magnitude history ' + r'($\nabla_{W1}$)')
ax2.set_ylim(0,1)
fig.tight_layout()
plt.show()
```

Final training loss: 0.10926794610680679



[5pts] Q2.4 Using ReLU activation

The Rectified Linear Unit (ReLU) activation is also widely used: $ReLU(x) = \max(0, x)$.

- Complete the implementation for the ReLU activation (forward and backward) in `mlp.py`.
- Train the network with ReLU, and report your final training loss.

Make sure you first pass the numerical gradient check on toy data.

In [6]:

```
net = init_toy_model('relu', std=1e-1)

loss, grads = net.loss(X, y, reg=0.1)
print('loss = ', loss) # correct_loss = 1.320973

# The differences should all be very small
print('checking gradients')
for param_name in grads:
    f = lambda W: net.loss(X, y, reg=0.1)[0]
    param_grad_num = eval_numerical_gradient(f, net.params[param_name], verbose=False)
    print('%s max relative error: %e'%(param_name, rel_error(param_grad_num, grads[param_name])))
```

```
loss = 1.3037878913298206
checking gradients
W2 max relative error: 3.440708e-09
b2 max relative error: 3.865091e-11
W1 max relative error: 3.561318e-09
b1 max relative error: 8.994864e-10
```

Now that it's working, let's train the network. Does the net get stronger learning signals (i.e. gradients) this time? Report your final training loss.

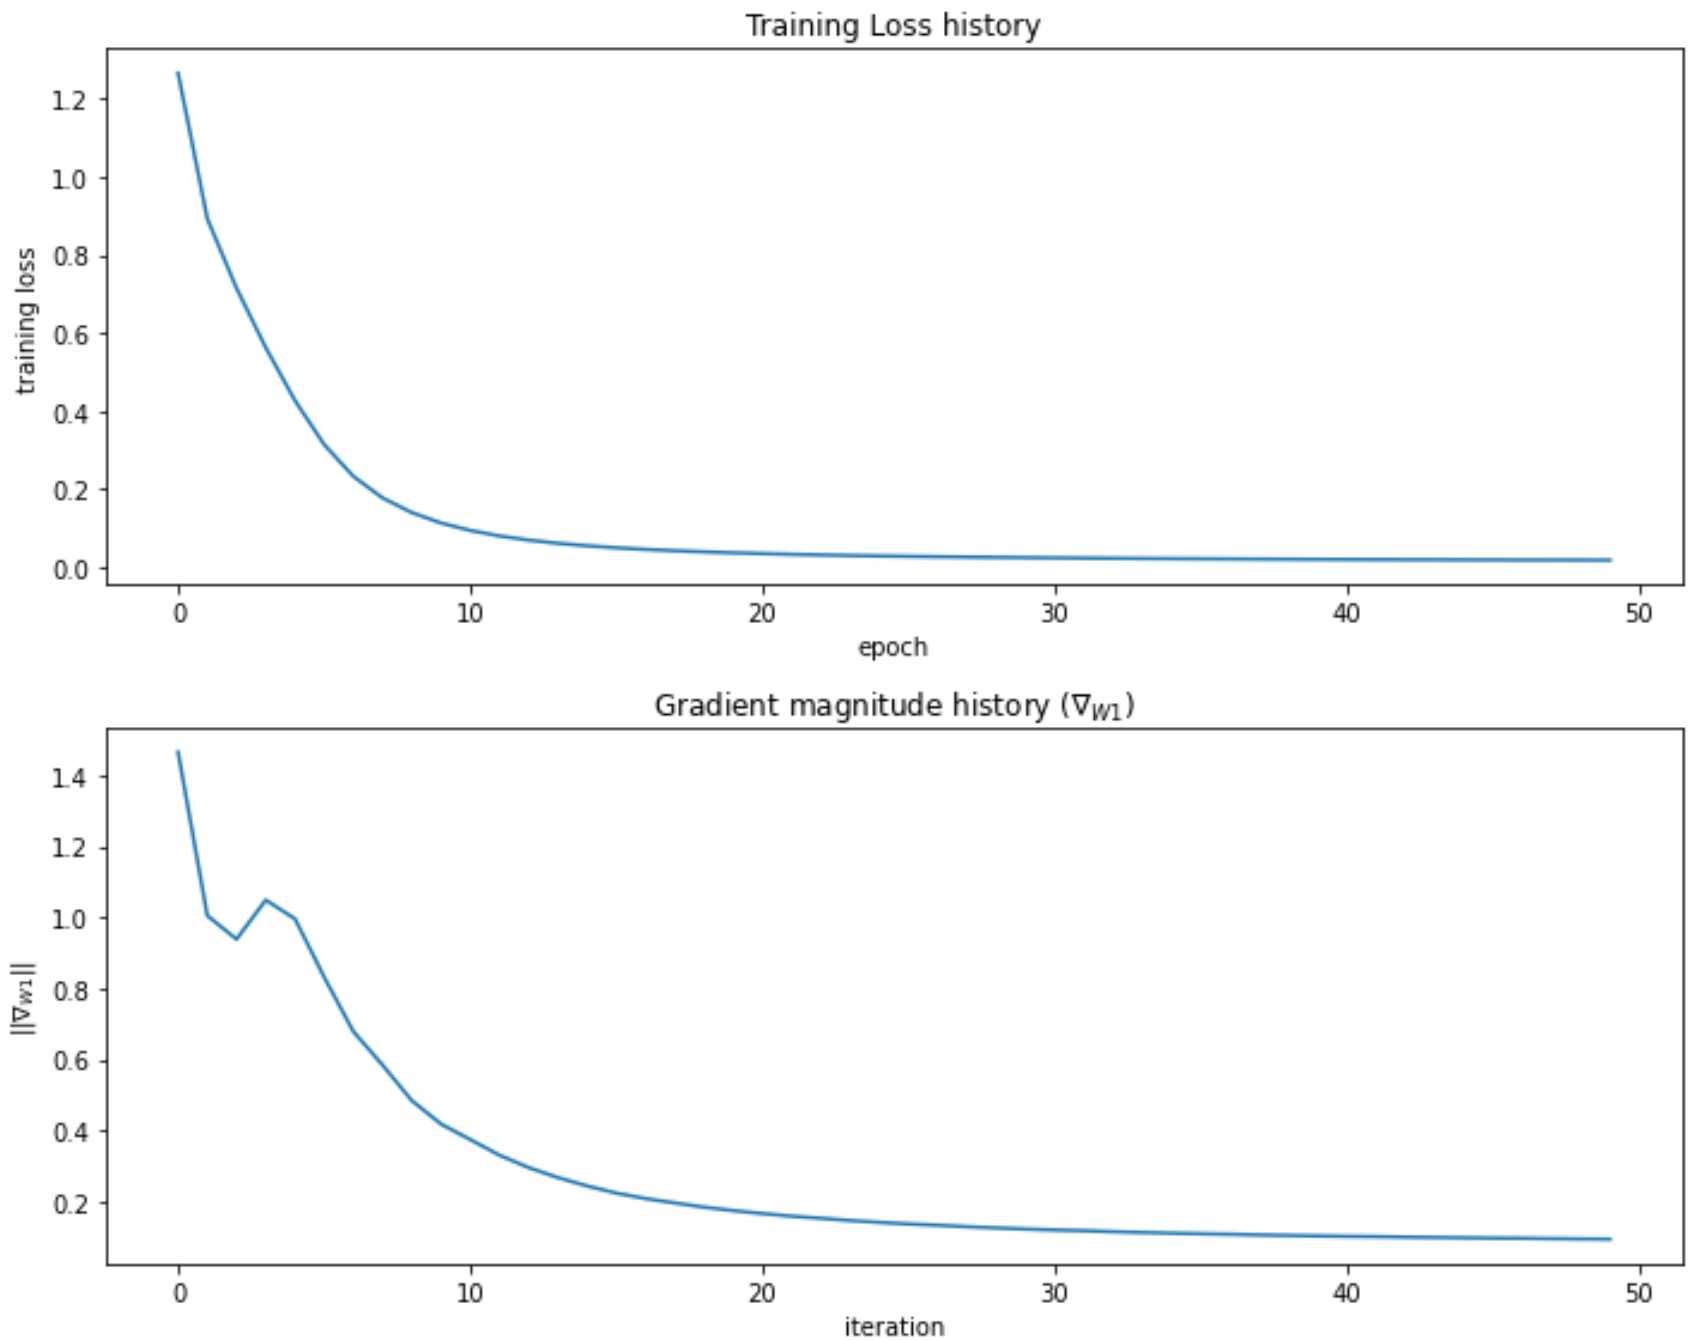
In [7]:

```
net = init_toy_model('relu', std=1e-1)
stats = net.train(X, y, X, y,
                  learning_rate=0.1, reg=1e-5,
                  num_epochs=50, verbose=False)

print('Final training loss: ', stats['loss_history'][-1])
```

```
# plot the loss history
fig, (ax1, ax2) = plt.subplots(2, 1)
ax1.plot(stats['loss_history'])
ax1.set_xlabel('epoch')
ax1.set_ylabel('training loss')
ax1.set_title('Training Loss history')
ax2.plot(stats['grad_magnitude_history'])
ax2.set_xlabel('iteration')
ax2.set_ylabel(r'$||\nabla_{W1}||$')
ax2.set_title('Gradient magnitude history ' + r'($\nabla_{W1}$)')
fig.tight_layout()
plt.show()
```

Final training loss: 0.0178562204869839



Application to a real Problem

Load MNIST data

Now that you have implemented a two-layer network that works on toy data, let's try some real data. The MNIST dataset is a standard machine learning benchmark. It consists of 70,000 grayscale handwritten digit images, which we split into 50,000 training, 10,000 validation and 10,000 testing. The images are of size 28x28, which are flattened into 784-d vectors.

Note 1: the function `get_MNIST_data` requires the `scikit-learn` package. If you previously did anaconda installation to set up your Python environment, you should already have it. Otherwise, you can install it following the instructions here: <http://scikit-learn.org/stable/install.html>

Note 2: If you encounter a HTTP 500 error, that is likely temporary, just try again.

Note 3: Ensure that the downloaded MNIST file is 55.4MB (smaller file-sizes could indicate an incomplete download - which is possible)

In [10]:

```
# Load MNIST
from utils import get_MNIST_data
X_train, y_train, X_val, y_val, X_test, y_test = get_MNIST_data()
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```

```
Train data shape: (50000, 784)
Train labels shape: (50000,)
Validation data shape: (10000, 784)
Validation labels shape: (10000,)
Test data shape: (10000, 784)
Test labels shape: (10000,)
```

Train a network on MNIST

We will now train a network on MNIST with 64 hidden units in the hidden layer. We train it using SGD, and decrease the learning rate with an exponential rate over time; this is achieved by multiplying the learning rate with a constant factor `learning_rate_decay` (which is less than 1) after each epoch. In effect, we are using a high learning rate initially, which is good for exploring the solution space, and using lower learning rates later to encourage convergence to a local minimum (or [saddle point](#), which may happen more often).

- Train your MNIST network with 2 different activation functions: sigmoid and ReLU.

We first define some variables and utility functions. The `plot_stats` function plots the histories of gradient magnitude, training loss, and accuracies on the training and validation sets. The `show_net_weights` function visualizes the weights learned in the first layer of the network. In most neural networks trained on visual data, the first layer weights typically show some visible structure when visualized. Both functions help you to diagnose the training process.

In [11]:

```
input_size = 28 * 28
hidden_size = 64
num_classes = 10

# Plot the loss function and train / validation accuracies
def plot_stats(stats):
    fig, (ax1, ax2, ax3) = plt.subplots(3, 1)
    ax1.plot(stats['grad_magnitude_history'])
    ax1.set_title('Gradient magnitude history ' + r'$(\nabla_{W1})$')
    ax1.set_xlabel('Iteration')
    ax1.set_ylabel(r'$||\nabla_{W1}||$')
    ax1.set_ylim(0, np.minimum(100, np.max(stats['grad_magnitude_history'])))

    ax2.plot(stats['loss_history'])
    ax2.set_title('Loss history')
    ax2.set_xlabel('Iteration')
    ax2.set_ylabel('Loss')
    ax2.set_ylim(0, 100)

    ax3.plot(stats['train_acc_history'], label='train')
```

```
ax3.plot(stats['val_acc_history'], label='val')
ax3.set_title('Classification accuracy history')
ax3.set_xlabel('Epoch')
ax3.set_ylabel('Classification accuracy')
fig.tight_layout()
plt.show()

# Visualize the weights of the network
from utils import visualize_grid
def show_net_weights(net):
    W1 = net.params['W1']
    W1 = W1.T.reshape(-1, 28, 28)
    plt.imshow(visualize_grid(W1, padding=3).astype('uint8'))
    plt.gca().axis('off')
    plt.show()
```

[10pts] Q2.5 Train a Sigmoid network

In [12]:

```
sigmoid_net = TwoLayerMLP(input_size, hidden_size, num_classes, activation='sigmoid', std=1e-1)

# Train the network
sigmoid_stats = sigmoid_net.train(X_train, y_train, X_val, y_val,
                                   num_epochs=20, batch_size=100,
                                   learning_rate=1e-3, learning_rate_decay=0.95,
                                   reg=0.5, verbose=True)

# Predict on the training set
train_acc = (sigmoid_net.predict(X_train) == y_train).mean()
print('Sigmoid final training accuracy: ', train_acc)

# Predict on the validation set
val_acc = (sigmoid_net.predict(X_val) == y_val).mean()
print('Sigmoid final validation accuracy: ', val_acc)

# Predict on the test set
test_acc = (sigmoid_net.predict(X_test) == y_test).mean()
```

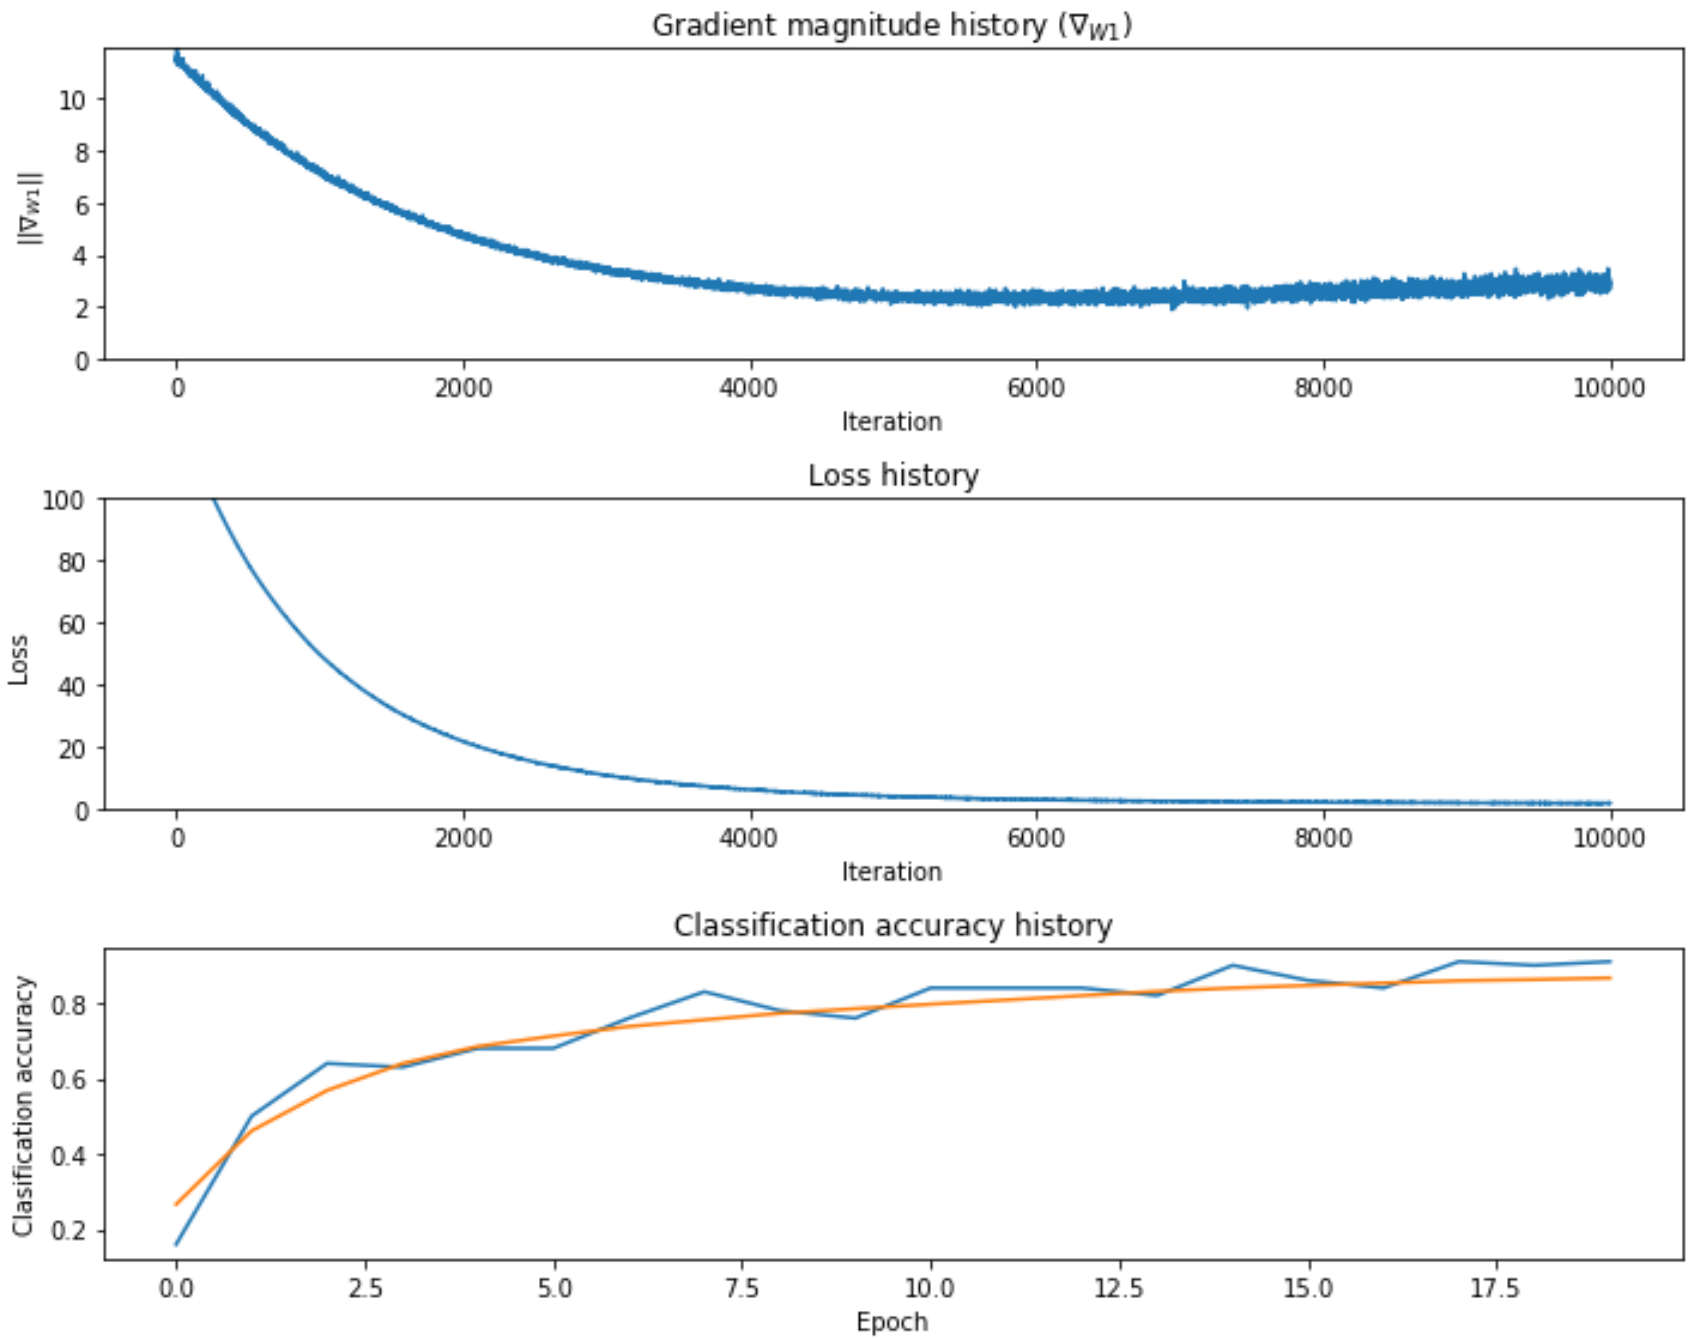


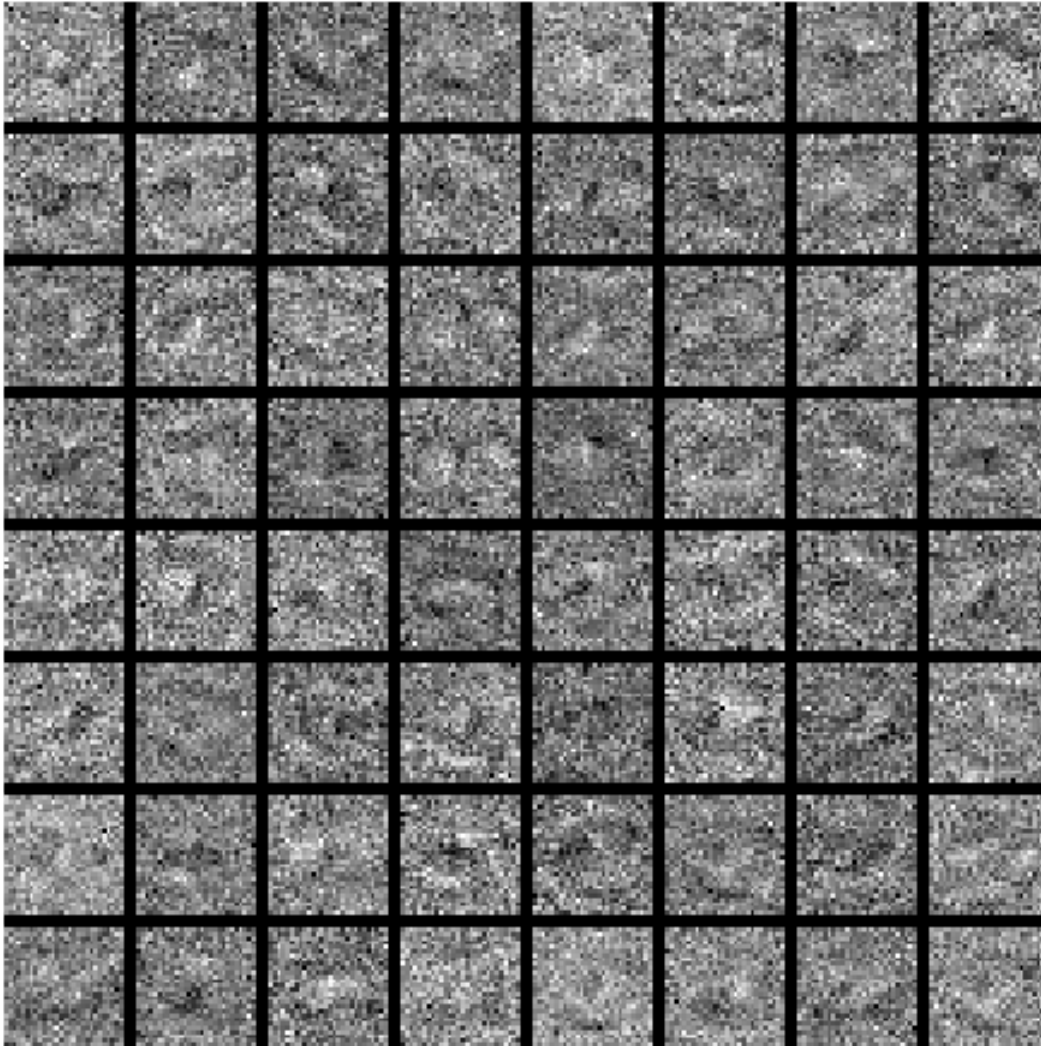
```
print('Sigmoid test accuracy: ', test_acc)

# show stats and visualizations
plot_stats(sigmoid_stats)
show_net_weights(sigmoid_net)
```

C:\Users\minhc\mlp.py:89: RuntimeWarning: overflow encountered in exp
hidden = 1/(1+np.exp(-z1))

```
Epoch 1: loss 79.040004, train_acc 0.160000, val_acc 0.266300
Epoch 2: loss 49.814996, train_acc 0.500000, val_acc 0.461100
Epoch 3: loss 32.419904, train_acc 0.640000, val_acc 0.568300
Epoch 4: loss 21.756599, train_acc 0.630000, val_acc 0.639700
Epoch 5: loss 15.148895, train_acc 0.680000, val_acc 0.685000
Epoch 6: loss 10.909900, train_acc 0.680000, val_acc 0.712600
Epoch 7: loss 8.078106, train_acc 0.760000, val_acc 0.737900
Epoch 8: loss 6.166522, train_acc 0.830000, val_acc 0.755600
Epoch 9: loss 4.948016, train_acc 0.780000, val_acc 0.772900
Epoch 10: loss 4.113118, train_acc 0.760000, val_acc 0.785000
Epoch 11: loss 3.455138, train_acc 0.840000, val_acc 0.797000
Epoch 12: loss 3.026239, train_acc 0.840000, val_acc 0.808100
Epoch 13: loss 2.702231, train_acc 0.840000, val_acc 0.819600
Epoch 14: loss 2.438965, train_acc 0.820000, val_acc 0.830900
Epoch 15: loss 2.258613, train_acc 0.900000, val_acc 0.839900
Epoch 16: loss 2.166625, train_acc 0.860000, val_acc 0.846800
Epoch 17: loss 2.098843, train_acc 0.840000, val_acc 0.852800
Epoch 18: loss 1.975990, train_acc 0.910000, val_acc 0.859300
Epoch 19: loss 1.898398, train_acc 0.900000, val_acc 0.862300
Epoch 20: loss 1.876564, train_acc 0.910000, val_acc 0.866400
Sigmoid final training accuracy: 0.8721
Sigmoid final validation accuracy: 0.8664
Sigmoid test accuracy: 0.8639
```





[10pts] Q2.6 Train a ReLU network

In [13]:

```
relu_net = TwoLayerMLP(input_size, hidden_size, num_classes, activation='relu', std=1e-1)

# Train the network
relu_stats = relu_net.train(X_train, y_train, X_val, y_val,
                             num_epochs=20, batch_size=100,
                             learning_rate=1e-3, learning_rate_decay=0.95,
```

```
reg=0.5, verbose=True)

# Predict on the training set
train_acc = (relu_net.predict(X_train) == y_train).mean()
print('ReLU final training accuracy: ', train_acc)

# Predict on the validation set
val_acc = (relu_net.predict(X_val) == y_val).mean()
print('ReLU final validation accuracy: ', val_acc)

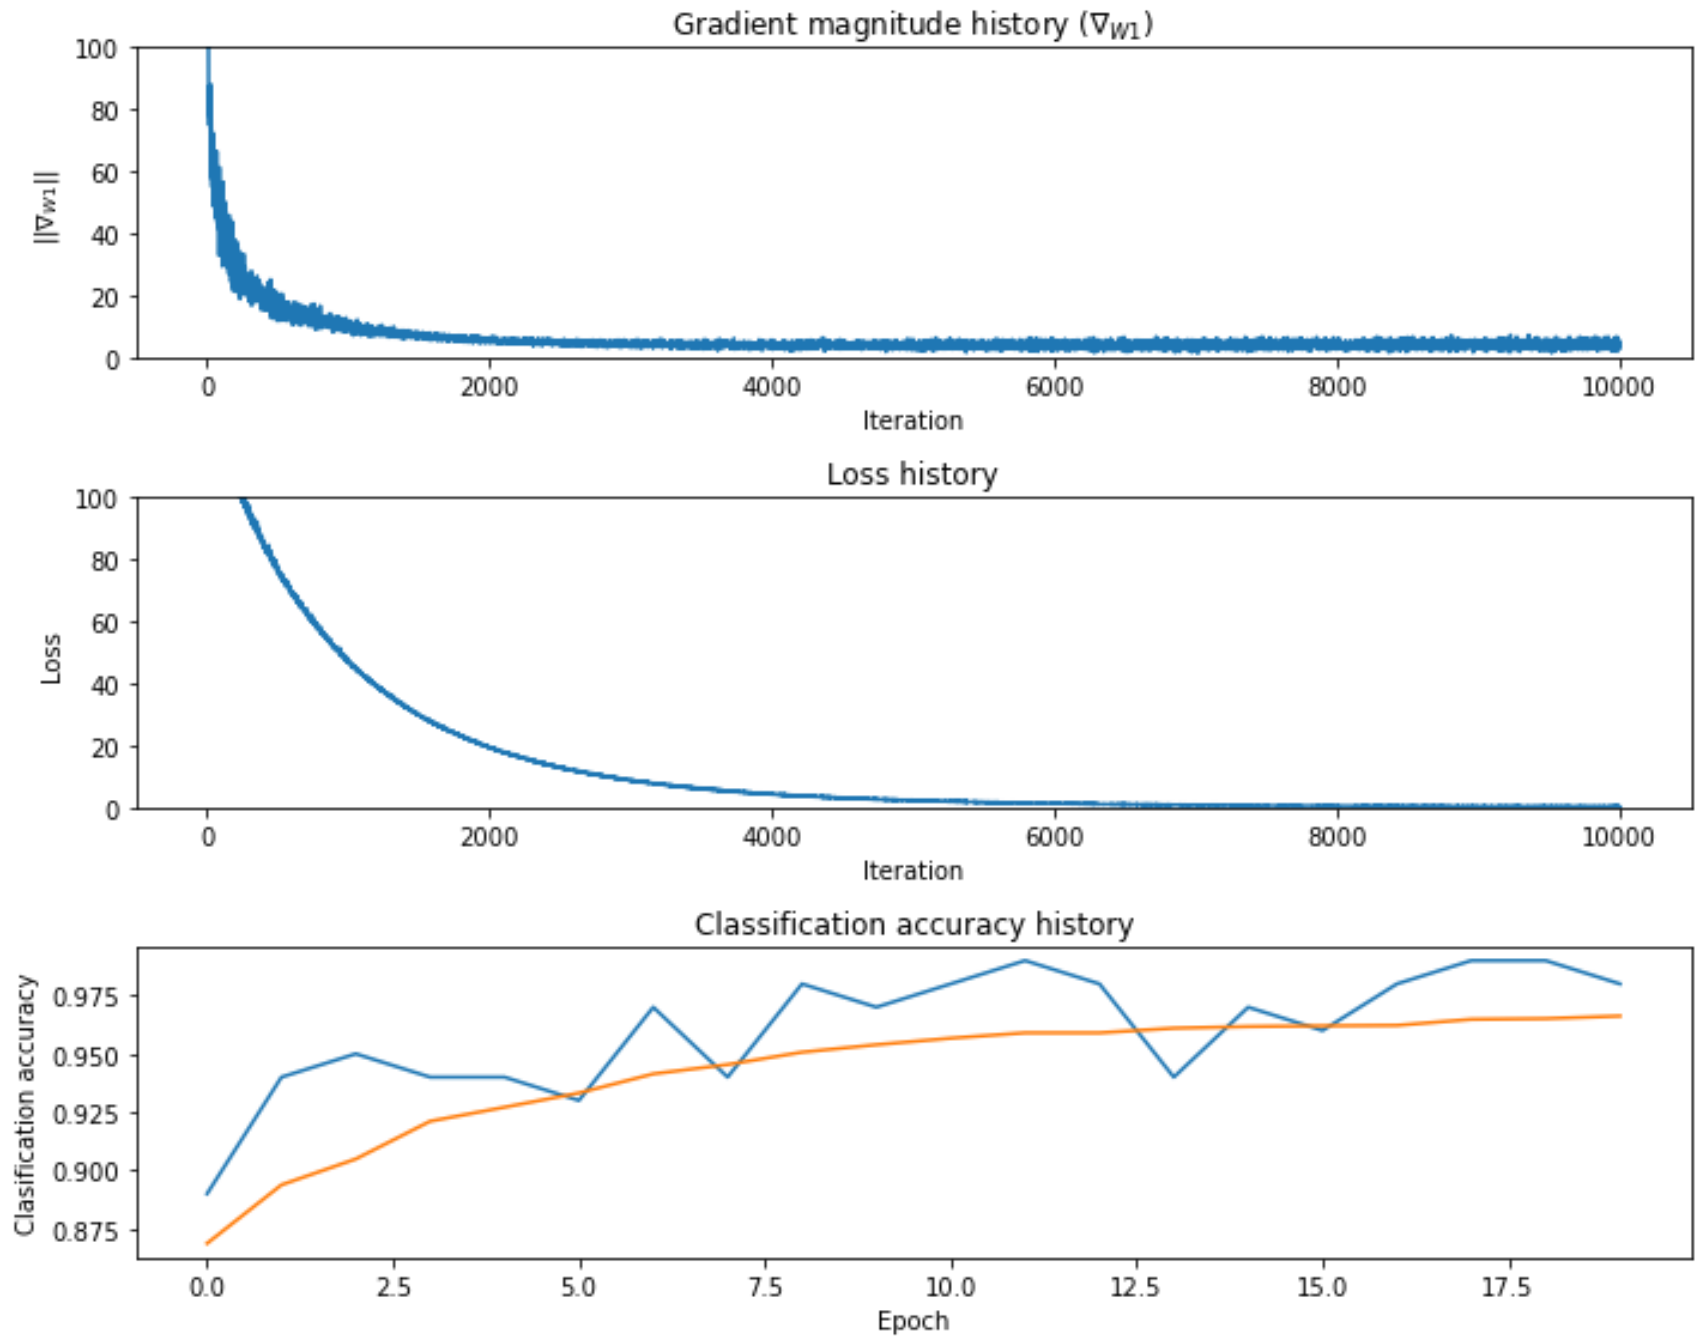
# Predict on the test set
test_acc = (relu_net.predict(X_test) == y_test).mean()
print('ReLU test accuracy: ', test_acc)

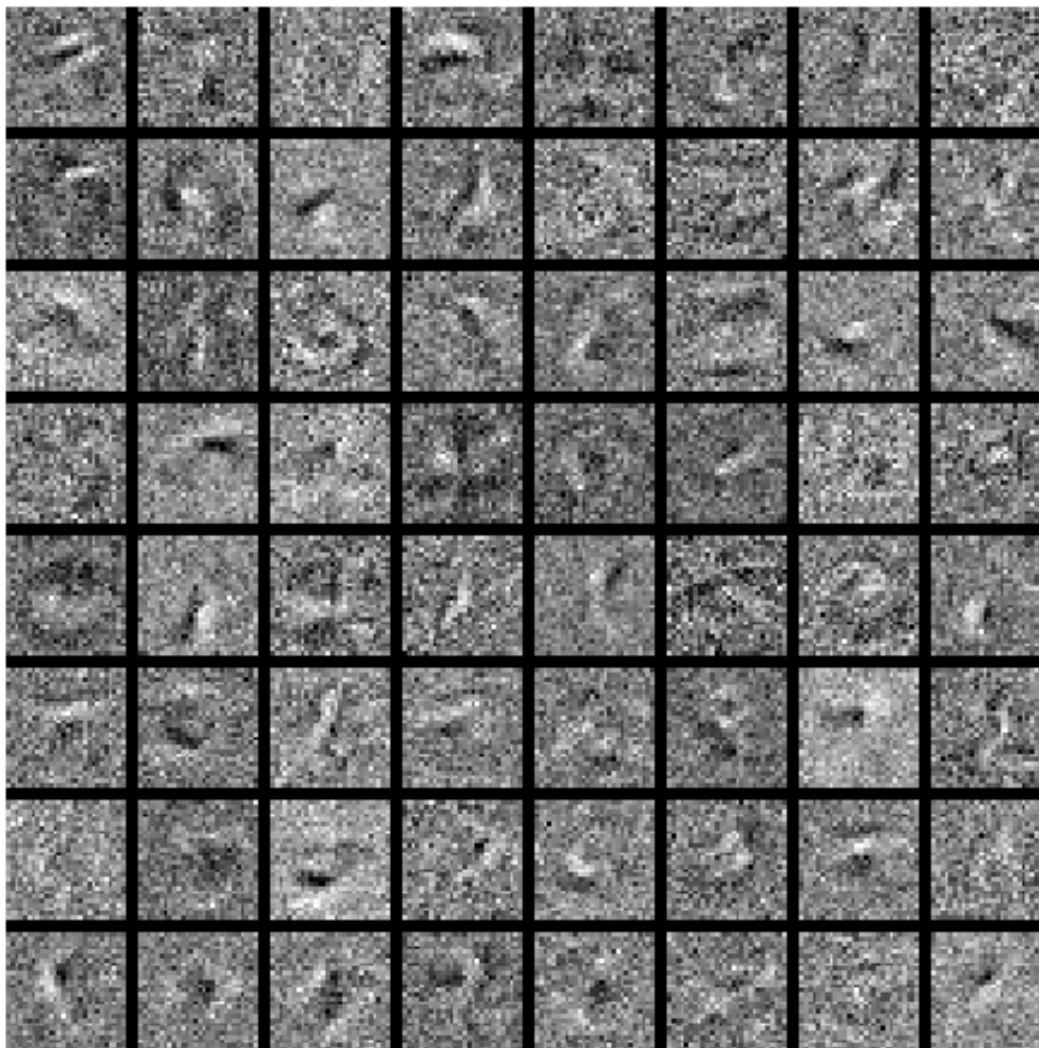
# show stats and visualizations
plot_stats(relu_stats)
show_net_weights(relu_net)
```

```
Epoch 1: loss 77.331905, train_acc 0.890000, val_acc 0.868700
Epoch 2: loss 47.510581, train_acc 0.940000, val_acc 0.893800
Epoch 3: loss 29.866540, train_acc 0.950000, val_acc 0.904900
Epoch 4: loss 19.523893, train_acc 0.940000, val_acc 0.921100
Epoch 5: loss 13.037804, train_acc 0.940000, val_acc 0.927100
Epoch 6: loss 8.990185, train_acc 0.930000, val_acc 0.933200
Epoch 7: loss 6.141586, train_acc 0.970000, val_acc 0.941400
Epoch 8: loss 4.525600, train_acc 0.940000, val_acc 0.945400
Epoch 9: loss 3.237260, train_acc 0.980000, val_acc 0.950700
Epoch 10: loss 2.411245, train_acc 0.970000, val_acc 0.953900
Epoch 11: loss 1.850464, train_acc 0.980000, val_acc 0.956700
Epoch 12: loss 1.394069, train_acc 0.990000, val_acc 0.959000
Epoch 13: loss 1.163258, train_acc 0.980000, val_acc 0.959000
Epoch 14: loss 0.994237, train_acc 0.940000, val_acc 0.961000
Epoch 15: loss 0.815998, train_acc 0.970000, val_acc 0.961700
Epoch 16: loss 0.682523, train_acc 0.960000, val_acc 0.962000
Epoch 17: loss 0.627242, train_acc 0.980000, val_acc 0.962200
Epoch 18: loss 0.488876, train_acc 0.990000, val_acc 0.964800
Epoch 19: loss 0.451647, train_acc 0.990000, val_acc 0.965100
Epoch 20: loss 0.404970, train_acc 0.980000, val_acc 0.966200
ReLU final training accuracy: 0.9729
```

ReLU final validation accuracy: 0.9662

ReLU test accuracy: 0.9655





[5pts] Q2.7

Which activation function would you choose in practice? Why?

ReLU has better accuracy comparing to Sigmoid. So I would choose ReLU.

[20pts] Problem 3: Simple Regularization Methods

You may have noticed the `reg` parameter in `TwoLayerMLP.loss`, controlling "regularization strength". In learning neural networks, aside from minimizing a loss function $\mathcal{L}(\theta)$ with respect to the network parameters θ , we usually explicitly or implicitly add some regularization term to reduce overfitting. A simple and popular regularization strategy is to penalize some *norm* of θ .

[10pts] Q3.1: L2 regularization

We can penalize the L2 norm of θ : we modify our objective function to be $\mathcal{L}(\theta) + \lambda \|\theta\|^2$ where λ is the weight of regularization. We will minimize this objective using gradient descent with step size η . Derive the update rule: at time $t + 1$, express the new parameters θ_{t+1} in terms of the old parameters θ_t , the gradient $g_t = \frac{\partial \mathcal{L}}{\partial \theta_t}$, η , and λ .

$$\begin{aligned}\theta_{t+1} &= \theta_t - \eta(\mathcal{L}(\theta) + \lambda \|\theta\|^2) \frac{\partial}{\partial \theta_t} \\ &= \theta_t - \eta\left(\frac{\partial \mathcal{L}(\theta)}{\partial \theta_t} + \frac{\partial \lambda \|\theta\|^2}{\partial \theta_t}\right) \\ &= \theta_t - \eta(g_t + 2\lambda\theta)\end{aligned}$$

[10pts] Q3.2: L1 regularization

Now let's consider L1 regularization: our objective in this case is $\mathcal{L}(\theta) + \lambda \|\theta\|_1$. Derive the update rule.

(Technically this becomes *Sub-Gradient* Descent since the L1 norm is not differentiable at 0. But practically it is usually not an issue.)

$$\theta_{t+1} = \theta_t - \eta(\mathcal{L}(\theta_t) + \lambda \|\theta_t\|_1) \frac{\partial}{\partial \theta_t}$$

$$= \theta_t - \eta(g_t + \lambda \sum_{i=1}^n |\theta_t^{(i)}| \frac{\partial}{\partial \theta_t})$$

$$= \theta_t - \eta(g_t + \lambda(-(\theta_t))) \text{ for } \theta_t < 0$$

$$\text{and } = \theta_t - \eta(g_t + \lambda(\theta_t)) \text{ for } \theta_t > 0$$

In []:

```
import numpy as np
import matplotlib.pyplot as plt

class TwoLayerMLP(object):
    """
    A two-layer fully-connected neural network. The net has an input dimension of
    N, a hidden layer dimension of H, and performs classification over C classes.
    We train the network with a softmax loss function and L2 regularization on the
    weight matrices. The network uses a ReLU nonlinearity after the first fully
    connected layer.

    In other words, the network has the following architecture:

    input - fully connected layer - ReLU - fully connected layer - softmax

    The outputs of the second fully-connected layer are the scores for each class.
    """

    def __init__(self, input_size, hidden_size, output_size, std=1e-4, activation='relu'):
        """
        Initialize the model. Weights are initialized to small random values and
        biases are initialized to zero. Weights and biases are stored in the
        variable self.params, which is a dictionary with the following keys:

        W1: First layer weights; has shape (D, H)
        b1: First layer biases; has shape (H,)
        W2: Second layer weights; has shape (H, C)
        b2: Second layer biases; has shape (C,)
        """
```


Inputs:

- input_size: The dimension D of the input data.
- hidden_size: The number of neurons H in the hidden layer.
- output_size: The number of classes C.

"""

```
self.params = {}
self.params['W1'] = std * np.random.randn(input_size, hidden_size)
self.params['b1'] = np.zeros(hidden_size)
self.params['W2'] = std * np.random.randn(hidden_size, output_size)
self.params['b2'] = np.zeros(output_size)
self.activation = activation
```

```
def loss(self, X, y=None, reg=0.0):
```

"""

Compute the loss and gradients for a two layer fully connected neural network.

Inputs:

- X: Input data of shape (N, D). Each X[i] is a training sample.
- y: Vector of training labels. y[i] is the label for X[i], and each y[i] is an integer in the range $0 \leq y[i] < C$. This parameter is optional; if it is not passed then we only return scores, and if it is passed then we instead return the loss and gradients.
- reg: Regularization strength.

Returns:

If y is None, return a matrix scores of shape (N, C) where scores[i, c] is the score for class c on input X[i].

If y is not None, instead return a tuple of:

- loss: Loss (data loss and regularization loss) for this batch of training samples.
- grads: Dictionary mapping parameter names to gradients of those parameters with respect to the loss function; has the same keys as self.params.

"""

Unpack variables from the params dictionary

```
W1, b1 = self.params['W1'], self.params['b1']
```

```

W2, b2 = self.params['W2'], self.params['b2']
_, C = W2.shape
N, D = X.shape

# Compute the forward pass
#####
# write your own code where you see [PLEASE IMPLEMENT]
#
# Perform the forward pass, computing the class scores for the input.
# Store the result in the scores variable, which should be an array of
# shape (N, C).
#####
z1 = np.dot(X, W1) + b1 # 1st layer activation, N*H

# 1st layer nonlinearity, N*H
if self.activation is 'relu':
    hidden = np.maximum(0, z1)
    # [PLEASE IMPLEMENT]
    #raise NotImplementedError('ReLU forward not implemented')
elif self.activation is 'sigmoid':
    # [PLEASE IMPLEMENT]
    hidden = 1/(1+np.exp(-z1))
    #raise NotImplementedError('Sigmoid forward not implemented')
else:
    raise ValueError('Unknown activation type')

# [PLEASE IMPLEMENT] 2nd layer activation, N*C
# hint: involves W2, b2
scores = np.dot(hidden, W2) + b2
#####
#                                     END OF YOUR CODE
#####

# If the targets are not given then jump out, we're done
if y is None:
    return scores

# cross-entropy loss with log-sum-exp

```

```

A = np.max(scores, axis=1) # N*1
F = np.exp(scores - A.reshape(N, 1)) # N*C
P = F / np.sum(F, axis=1).reshape(N, 1) # N*C
loss = np.mean(-np.choose(y, scores.T) + np.log(np.sum(F, axis=1)) + A)
# add regularization terms
loss += 0.5 * reg * np.sum(W1 * W1)
loss += 0.5 * reg * np.sum(W2 * W2)

# Backward pass: compute gradients
grads = {}
#####
# write your own code where you see [PLEASE IMPLEMENT]
#
# Compute the backward pass, computing the derivatives of the weights
# and biases. Store the results in the grads dictionary. For example,
# grads['W1'] should store the gradient on W1, and be a matrix of same size.
# You should define the hidden variable in the part where you implement the
# different activation functions. "hidden" is the output after you apply
# the activation function for the hidden layer.
# Hint: you should apply different activation functions on "z1" and get "hidden".
#####

# output layer
yc = np.zeros((N,C))
for i in range(N):
    yc[i, y[i]] = 1

dscore = P - yc # [PLEASE IMPLEMENT] partial derivative of loss wrt. the logits (dL/dz)
dW2 = np.dot(hidden.T, dscore)/N # partial derivative of loss wrt. W2
db2 = np.mean(dscore, axis=0) # partial derivation of loss wrt. b2

# hidden layer
dhidden = np.dot(dscore, W2.T)
if self.activation is 'relu':
    # [PLEASE IMPLEMENT]
    dz1 = dhidden
    dz1[z1<0] = 0

```

```

        #raise NotImplementedError('ReLU backward not implemented')
    elif self.activation is 'sigmoid':
        # [PLEASE IMPLEMENT]
        dz1 = (hidden*(1-hidden))*dhidden
        #raise NotImplementedError('Sigmoid backward not implemented')
    else:
        raise ValueError('Unknown activation type')

    # first layer
    dW1 = np.dot(X.T, dz1)/N # [PLEASE IMPLEMENT]
    db1 = np.mean(dz1, axis = 0) # [PLEASE IMPLEMENT]
    #####
    #                                END OF YOUR CODE
    #####

    grads['W2'] = dW2 + reg*W2
    grads['b2'] = db2
    grads['W1'] = dW1 + reg*W1
    grads['b1'] = db1
    return loss, grads

def train(self, X, y, X_val, y_val,
          learning_rate=1e-3, learning_rate_decay=0.95,
          reg=1e-5, num_epochs=10,
          batch_size=200, verbose=False):
    """
    Train this neural network using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) giving training data.
    - y: A numpy array of shape (N,) giving training labels; y[i] = c means that
      X[i] has label c, where 0 <= c < C.
    - X_val: A numpy array of shape (N_val, D) giving validation data.
    - y_val: A numpy array of shape (N_val,) giving validation labels.
    - learning_rate: Scalar giving learning rate for optimization.
    - learning_rate_decay: Scalar giving factor used to decay the learning rate

```

```
    after each epoch.
- reg: Scalar giving regularization strength.
- num_iters: Number of steps to take when optimizing.
- batch_size: Number of training examples to use per step.
- verbose: boolean; if true print progress during optimization.
"""

num_train = X.shape[0]
iterations_per_epoch = max(int(num_train / batch_size), 1)
epoch_num = 0

# Use SGD to optimize the parameters in self.model
loss_history = []
grad_magnitude_history = []
train_acc_history = []
val_acc_history = []

np.random.seed(1)
for epoch in range(num_epochs):
    # fixed permutation (within this epoch) of training data
    perm = np.random.permutation(num_train)

    # go through minibatches
    for it in range(iterations_per_epoch):
        X_batch = None
        y_batch = None

        # Create a random minibatch
        idx = perm[it*batch_size:(it+1)*batch_size]
        X_batch = X[idx, :]
        y_batch = y[idx]

        # Compute loss and gradients using the current minibatch
        loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
        loss_history.append(loss)

        # do gradient descent
        for param in self.params:
            self.params[param] -= grads[param] * learning_rate
```

```

        # record gradient magnitude (Frobenius) for W1
        grad_magnitude_history.append(np.linalg.norm(grads['W1']))

    # Every epoch, check train and val accuracy and decay Learning rate.
    # Check accuracy
    train_acc = (self.predict(X_batch) == y_batch).mean()
    val_acc = (self.predict(X_val) == y_val).mean()
    train_acc_history.append(train_acc)
    val_acc_history.append(val_acc)
    if verbose:
        print('Epoch %d: loss %f, train_acc %f, val_acc %f'%(
            epoch+1, loss, train_acc, val_acc))

    # Decay Learning rate
    learning_rate *= learning_rate_decay

return {
    'loss_history': loss_history,
    'grad_magnitude_history': grad_magnitude_history,
    'train_acc_history': train_acc_history,
    'val_acc_history': val_acc_history,
}

def predict(self, X):
    """
    Use the trained weights of this two-layer network to predict labels for
    data points. For each data point we predict scores for each of the C
    classes, and assign each data point to the class with the highest score.

    Inputs:
    - X: A numpy array of shape (N, D) giving N D-dimensional data points to
        classify.

    Returns:
    - y_pred: A numpy array of shape (N,) giving predicted labels for each of
        the elements of X. For all i, y_pred[i] = c means that X[i] is predicted

```

```
    to have class c, where  $0 \leq c < C$ .  
    ""
```

```
#####  
# [PLEASE IMPLEMENT]  
# hint: it should be very easy  
y_pred = np.argmax(self.loss(X), axis = 1)  
#####  
  
return y_pred
```